

Linux capable RISC-V CPU for IOb-SoC

Pedro Nuno de Melo Antunes

Thesis to obtain the Master of Science Degree in

Electrical and Computer Engineering

Supervisor(s): Prof. José João Henriques Teixeira de Sousa

Examination Committee

Supervisor: Prof. José João Henriques Teixeira de Sousa

Member of the Committee: Prof. Horácio Cláudio De Campos Neto

October 7, 2022

Abstract

Blablabla

Honor Declaration

“Declaro por minha honra de que ”

Local, Date

Signature

Contents

List of Figures	I
List of Tables	III
List of Code	V
List of Acronyms	VII
1 Introduction	1
1.1 Motivation	1
1.2 Objectives	1
2 Must Have Concepts	3
2.1 The <i>IOb-SoC</i> template	3
2.1.1 Adding peripherals	4
2.1.2 Internal Buses	4
2.1.3 <i>iob-split</i> and <i>iob-merge</i>	4
2.1.4 Bootloader and bare-metal firmware	5
2.2 RISC-V	5
2.2.1 CLINT Specification	5
2.2.2 PLIC Specification	5
2.2.3 UART/Serial Console	5
2.3 Open Source Verification tools	6
2.3.1 Qemu Simulation	6
2.4 The Linux Boot Flow on RISC-V	6
2.4.1 Bootloader firmware	6
2.4.2 What is a device tree?	6
3 Existing Embedded Technologies	7
3.1 Closed source RISC-V Embedded Systems	7
3.1.1 Andes Technology	8
3.1.2 SiFive	8
3.2 Open-Source Solutions	9
3.2.1 <i>CVA6</i>	9
3.2.2 The Berkeley Out-of-Order RISC-V Processor	10
3.2.3 <i>VexRiscv</i>	11
3.3 Overall CPU comparison	12
4 Hardware Developed	15
4.1 Central Processing Unit	16
4.1.1 VexRiscv Wrapper	18

4.2	<i>UART 16550</i>	22
4.2.1	<i>UART 16550 Wrapper</i>	22
4.3	CLINT Unit	24
4.4	PLIC Unit	25
4.4.1	PLIC wrapper	26
5	Software Developed	29
5.1	<i>Python Console</i>	29
5.2	<i>IOb-SoC Simulation</i>	31
5.2.1	UUT Top Hardware module	31
5.2.2	Verilator Testbench	31
5.3	Interrupt Routine	32
5.3.1	CLINT simulation	34
5.3.2	Bare-metal firmware	34
5.4	<i>IOb-SoC Linux OS integration</i>	35
5.4.1	Bootloaders	35
5.4.2	Device Tree	36
5.4.3	Linux kernel	36
5.4.4	Root File system	36
5.5	Makefiles	36
6	Project Results	37
6.1	System Running "Hello World!"	37
6.2	Interrupt Routines	37
6.3	Run/Boot Linux Performance	37
6.4	FPGA Resources Consumption	37
7	Conclusions	39
7.1	Achievements	39
7.2	Contributed Repositories	39
7.3	Future Work	39
A	Annex	VII
A.1	Annex 1	VII
A.2	Annex 2	VII
	Bibliography	IX

List of Figures

2.1	<i>IOb-SoC</i> sketch.	4
3.1	CVA6 core design architecture.	10
4.1	Developed SoC sketch.	15
5.1	Simulated hardware interfaces.	32

List of Tables

3.1	CPU comparison table: Y means the CPU supports the feature; X means the CPU does not support the feature; N/A means the feature does not apply to the respective CPU.	13
4.1	<i>VexRiscv</i> core inputs and outputs.	20
4.2	First try at identifying the rules the <code>cmd_ready</code> should follow.	21
4.3	Simplified truth table.	21
4.4	WISHBONE interface signals.	23
4.5	<i>UART16550</i> interface with <i>IOb-SoC</i>	23
4.6	CLINT interface with <i>IOb-SoC</i>	24
4.7	PLIC interface with <i>IOb-SoC</i>	26

List of Code

4.1	Generate <i>verilog</i> from <i>SpinalHDL</i>	17
5.1	Call <i>Console</i> program	31
5.2	Set Up Timer Interrupt.	33
5.3	Read core id from CSR.	33
5.4	Set Up Software Interrupt.	33

List of Acronyms

ACK Acknowledgement.

ALU arithmetic logic unit.

AMO Atomic Memory Operations.

ASCII American Standard Code for Information Interchange.

ASIC Application-Specific Integrated Circuit.

Chisel Constructing Hardware in a Scala Embedded Language.

CLINT Core-local Interrupt Controller.

CPU Central Processing Unit.

CSR Control and status register.

DC1 Device Control 1.

DEMUX demultiplexer.

ENQ Enquiry.

EOT End of Transmission.

FPGA Field-programmable gate array.

HDL Hardware Description Language.

ISA Instruction set architecture.

LR Load-Reserved.

lsb Less Significant bit.

M Machine.

MMU Memory Management Unit.

msb Most Significant bit.

MUX multiplexer.

OoO Out-of-Order.

List of Acronyms

OS Operating System.

PC Program counter.

PLIC Platform-Level Interrupt Controller.

rtc arithmetic logic unit.

RTL register-transfer level.

S Supervisor.

SC Store-Conditional.

SoC System on a chip.

U User.

UART Universal asynchronous receiver/transmitter.

UUT Unit Under Test.

1 | Introduction

1.1 Motivation

1.2 Objectives

2 | Must Have Concepts

During this chapter, it is going to be discuss topics that help understand the technology developments made along the thesis project. The developments will involve both hardware and software components. As such there are hardware and software concepts that are important to have before discussing the next chapters.

In this project a SoC is going to be developed. But it is not going to be created from scratch. The *IOb-SoC* is going to be used as a starting point. Consequently it is important to understand how it works beforehand. It is also important to study the RISC-V Instruction set architecture (ISA). Since the hardware developed in this project will be compatible with the RISC-V ISA. Additionally, the software and firmware used will be compiled with the RISC-V cross-compiler toolchain. An important part when developing a system is its testing and simulation before implementation. Therefor, an overview of available methods for simulation of the developed components is going to be seen. Finally, the last concept that is important to this project is the boot flow of an Operating System (OS) on a RISC-V platform.

2.1 The *IOb-SoC* template

The *IOb-SoC* [10] is a System-on-Chip (SoC) template that eases the creation of a new SoC. The *IOb-SoC* provides a base Verilog hardware design equipped with an open-source RISC-V processor, an internal SRAM memory subsystem, a UART, and an optional interface to external memory. If the external memory interface is selected, an instruction L1 cache, a data L1 cache and a shared L2 cache are added. The L2 cache communicates with a third party memory controller IP (typically a DDR controller) using an AXI4 master bus. Users can add IP cores and software to build their own SoCs quickly. This way, hardware accelerators can be easily created and tested with the developed firmware.

In figure 2.1 it is represented a sketch of the SoC design. This design is valid at the start of this project. During the hardware developed chapter 4 some alterations were made to the *IOb-SoC* original template.

Building a new processor-based system from scratch can be difficult. *IOb-SoC* has been created to facilitate this process. In this work, a variant of the existing *IOb-SoC* capable of running a Linux Operating System is developed. *IOb-SoC* currently supports three FPGA board models: the Xilinx Kintex UltraScale KU040 Development Board, the Basys 3 Artix-7 FPGA Trainer Board and the Cyclone V GT FPGA Development Kit.

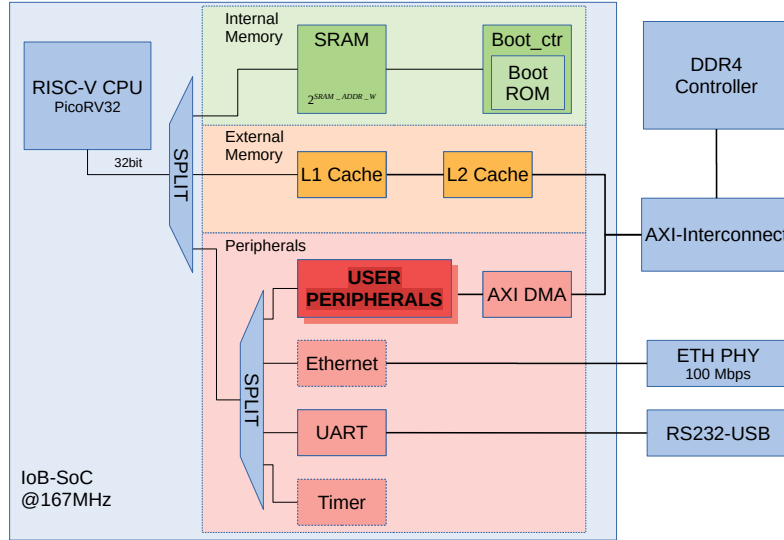


Figure 2.1: IOb-SoC sketch.

2.1.1 Adding peripherals

2.1.2 Internal Buses

Review the “cpu_i_resp”, “cpu_d_resp”, “cpu_i_req” and “cpu_d_req” signals.

2.1.3 *iob-split* and *iob-merge*

The ***iob-split*** is simply a configurable desmultiplexer (DEMUX). Meaning that when the *iob-split* hardware module is instantiated the developer can configure it. The developer is able to change the size of the desmultiplexer and the selection bits, through N_SLAVES and P_SLAVES respectively. N_SLAVES corresponds to the number of slaves, witch can also be seen as the number of the DEMUX outputs. P_SLAVES indicates the slave select word Most Significant bit (msb) position, in other words it is the position of the msb of the desmultiplexer selection bits. The number of the selection bits is given by equation 2.1.

$$Nb = \log_2(N_SLAVES) + (\log_2(N_SLAVES) == 0) \quad (2.1)$$

The ***iob-merge*** works similar to the *iob-split* but instead of being a DEMUX it is a configurable multiplexer (MUX). Meaning that instead of having multiple outputs and one input it has multiple inputs and one output. The number of inputs is indicated by N_SLAVES and the selection bits are chosen by P_SLAVES.

2.1.4 Bootloader and bare-metal firmware

2.2 RISC-V

Talk about the 32 registers in the register file

The instruction each ISA extension contains...

Control and status register (CSR) needed to run a full feature OS... (`core_id`, `misa`, `mcause`, ...)

2.2.1 CLINT Specification

The RISC-V **CLINT** is described ... Platform must support an ACLINT MTIME counter resolution of 100ns or less (corresponding to a clock tick frequency of at least 10 MHz).

2.2.2 PLIC Specification

The RISC-V **PLIC** was first described in the privilege instructions documentation, but since version 1.10 it was moved to its own document.

2.2.3 UART/Serial Console

In the RISC-V Platform Specification [18] it is defined that every embedded Operating System (OS) is required to have a UART port implementation that is register-compatible with the industry standard *UART 16550*, which was studied in chapter ???. The *UART 16550* already exists for a long time, it was released by *National Semiconductor* in 1987, and is present and supported by a large number of software and hardware. The *UART 16550* is often used connected to an RS-232 interface and in this project the development boards used are connected through RS-232 to the computer.

The *UART 16550* registers are...

2.3 Open Source Verification tools

For testing purposes, it is important to have a good hardware simulation environment. For that, we take advantage of already existing and well-developed tools. There exist a number of simulation tools, most of them are proprietary, as for example 'xcelium' from 'Candence'. Its utilization can increase the cost of a project significantly. In this Thesis we will make an effort of using open-source, free to use, verification tools. In specific, we will take advantage of 'Icarus Verilog' and 'Verilator'. Although both tools are for verification, they serve different purposes, due to their characteristics.

- '**Icarus Verilog**' is a Logic Simulator that uses verilog or system-verilog testbenches to test the UUT (Unit Under Test). Unfortunately, its support for system-verilog is limited and some designs might not run in this simulator. '*Icarus Verilog*' is also known as '*IVerilog*'.

After compiling the hardware design, '*IVerilog*' outputs a file which can be run line by line to simulate designed logic.

- '**Verilator**'

The biggest differences are: '*Verilator*' only represents logic signal as 1's or 0's, contrary to '*IVerilog*' which also represents unknown values as X's; Since '*Verilator*' ends up being a C++ program it is much faster to run the simulation than with '*IVerilog*'; On another perspective '*Verilator*' is slower than '*IVerilog*' to interpret the hardware logic design. As such, it is easier to use '*IVerilog*' to detect errors on the design, but it is better to use '*Verilator*' for more complexed simulations.

2.3.1 Qemu Simulation

<https://www.sifive.com/blog/risc-v-qemu-part-1-privileged-isa-hifive1-virtio>

2.4 The Linux Boot Flow on RISC-V

2.4.1 Bootloader firmware

2.4.2 What is a device tree?

3 | Existing Embedded Technologies

There already exists embedded microcontrollers capable of running Linux. Big companies for example ARM, Qualcomm, MediaTek, Intel and AMD have created microcontrollers capable of running Linux. But the processor architecture of those microcontrollers is not open-source, much less the microcontroller itself.

As an example, the *Raspberry Pi 4* is a very capable and cheap board where a developer can test and implement new software running in Linux. The Raspberry CPU is an *Cortex-A72* [1] which is a System on Chip (SoC) developed by ARM on their ARMv8 64-bit CPU architecture. But if someone wanted to use the Raspberry as a base for his costume hardware design, that would be impossible. And thus appears the need for open-source hardware that allows creating something new without having to start from scratch every time. This led to the appearance of RISC-V the open-source CPU architecture.

3.1 Closed source RISC-V Embedded Systems

Since then, a few companies using RISC-V have appeared. RISC-V CPUs are already present in the automotive and IoT markets, besides AI chips in data centres. Due to the RISC-V ISA royalty-free license, new StartUps tend to look at RISC-V CPUs as a solution for their cores. Even if the CPU Core isn't free to use it ends up being a cheaper solution.

While creating new products companies proved how advantageous the RISC-V architecture was. Furthermore, they have contributed to open-source software, hardware and documentation. Some companies with big recognition involved with RISC-V technology are:

- *Western Digital* who now uses RISC-V in its external storage disks.
- *Microchip* as launched the first RISC-V-Based System-on-Chip (SoC) FPGA, *PolarFire*.
- *Antmicro/Microsemi* ¹ have built a software called Renode that is used to develop, debug and test multi-node RISC-V device systems.
- *BeagleBoard.org*, *Seeed Studio* and *StarFive* worked together to build the first affordable RISC-V computer designed to run Linux, *BeagleV* [5]. The board is priced around 150€.

These companies have all helped pave the way for a full-feature Operating System based on the Linux kernel to be compatible with the RISC-V architecture. However, two companies have a bigger impact on RISC-V CPU design, those are Andes Technology and SiFive.

¹Microchip has acquired Microsemi Corporation in May 2018.

3.1.1 Andes Technology

Andes Technology is one of the founding members of RISC-V International. Since it is highly involved with RISC-V it ended up being one of the major contributors (and maintainers) of the RISC-V toolchain. This is important because the RISC-V ISA is merely an instruction set architecture, there needs to exist complementing software, such as compiler and development tools.

Nowadays Andes CPUs are applied nearly everywhere, from telecommunications, storage controllers, and touch screen sensors to data centres, etc. Andes Technologies has had incredible success using RISC-V technology, as proof they have shipped billions of embedded SoC with RISC-V processors based on their RISC-V ISA variant, AndeStar™ V5.

Andes CPUs which are capable of running Linux are the *A25* [29] and *AX25* [30]. Both support single and double precision floating points, the RISC-V P-extension (draft) DSP/SIMD ISA and an MMU (Memory Management Unit) for Linux applications. Besides that both enable the use of Machine (M), User (U) and Supervisor (S) Privilege levels that allow running Linux and other advanced operating systems with protection between kernel and user programs. Furthermore, both have L1 instructions and data cache. The difference between them is that *A25* is based on 32-bit architecture and the *AX25* is 64-bit. This leads to the *AX25* being ideal for embedded applications that need to access address space over 4GB, and the *A25* being smaller in gate count. Both CPUs can be implemented on the *AE350* [28] SoC allowing to use these CPUs on developer boards, for example in the *ADP-XC7K160/410* [27].

3.1.2 SiFive

SiFive is a company that was born from the RISC-V ISA. SiFive was founded by three researchers from the University of California Berkeley, Krste Asanović, Yunsup Lee, and Andrew Waterman. Those researchers were deeply involved with the development of the RISC-V ISA, from working on the base ISA to working on the floating point numbers and compressed instructions ISA extensions. It is no surprise that the first company to release a chip and development board that implemented the RISC-V ISA was SiFive. This happened in 2016 one year after the company was founded.

In 2017 SiFive launched *U54* [22] which was the first RISC-V CPU capable of running a full fledged Operating System like Linux. With it they launched the *U54-MC* [23] SoC that had four *U54* 64-bit cores. Furthermore, the *U54-MC* implemented the initial CLINT and PLIC unit. The development of the CLINT and the PLIC made by SiFive would eventually lead to the documentation and specification of the respective hardware components with which RISC-V systems must be compliant if they proclaim to use either one. One year after, in 2018, they launched *HiFive Unleashed* [20] which was the first board that implemented the *U54* CPU and

run a Linux OS with a desktop environment (DE). The *HiFive Unleashed* has been discontinued and better hardware has been made available.

SiFive has since then extended its *U Cores* product lineup. All *U* cores are 64-bit application processors capable of running Linux. The highest performance core is the *U74* [24]. The core architecture is RV64GBC which means it supports the RISC-V I, M, A, F, D, B and C ISA extensions (explained in ***ref section 2.3.x***). This CPU has already been applied to multiple boards, for example, the *BeagleV* has a SoC with dual-core SiFive *U74* CPU. SiFive also launched its own development board, *HiFive Unmatched* [21], with four *U74* cores on the *U74-MC* [25] SoC. Furthermore, in 2021, Canonical the developers behind Ubuntu announced the OS support for both the HiFive Unmatched and HiFive Unleashed.

3.2 Open-Source Solutions

Built upon the RISC-V open-source Instruction set architecture, various CPU designs have emerged. Some of them are fully open-source and might be implemented in other projects. Those CPUs were mostly developed by Universities research groups or by individuals with a grant.

RISC-V CPUs are most popular in embedded systems and IoT devices. Consequently there exists a wide variety of open-source CPUs which are implemented on multiple embedded microcontrollers. A few examples of those CPUs would be the *PicoRV32* [31], *NEORV32* [26], *DarkRISCv* [7] and *Ibex* [12] from lowRISC. But those will not be discussed in detail in this paper since they do not meet the requirements to run the Linux Kernel. These CPUs either only support Machine (M) level privilege mode or support Machine (M)+Supervisor (S) mode. Moreover, none of the given examples supports the Atomic RISC-V ISA extension. This extension is essential to run Linux. Since the kernel explicitly executes instructions from the Atomic extension.

To run a Linux-based Operating System an application processor is needed. A CPU is considered an *application processor* if it has the hardware required to run a full-feature Operating System (OS) and user applications. This means that the processor should have the required Control and status register (CSR), support M+S+U privilege modes and support atomic instructions. An open-source solution would be either the *CVA6* [32] (previously known as Ariane), *BOOM* [33] or *VexRiscv* [16].

3.2.1 CVA6

The CVA6 is a 6-stage, single issue, in-order CPU which can execute either the 32-bit or 64-bit RISC-V instruction set. CVA6 has support for the I, M, A and C RISC-V ISA extensions. The original design was initiated in a research group by a PhD student at ETH Zurich (where

they called the core Ariane). Since then the development and maintenance of CVA6 were incorporated in the *OpenHW Group* as part of their CORE-V processor lineup. The support for RV32IMAC was only developed recently by Thales and is also open-source. The CPU design is illustrated in figure 3.1 that was obtained from: <https://github.com/openhwgroup/cva6/>.




Figure 3.1: CVA6 core design architecture.

The CVA6 supports any operating system based on Unix since it implements the three needed privilege levels M, S and U. The core is written in SystemVerilog, and its micro-architecture is designed to reduce the critical path length. Since it is written in SystemVerilog, it is easier for someone knowledgeable in the classic Verilog and VHDL languages to understand and create a customized CPU based on the CVA6, than if it was written in a high-level hardware description language. However, although the CVA6 is an open-source project it is hard to take advantage of isolated hardware components. This is a consequence of how it was developed. Every SystemVerilog module of the CVA6 depends on other files from the project and the CPU itself is very little customizable. To illustrate the problem if we wanted to remove the L1 cache present in the CVA6 to use the L1 cache used on IOB-SoC, it would be very difficult and time-consuming to create a CPU core without that component.

The CVA6 can be found implemented on *OpenPiton* [4]. *OpenPiton* is an open-source project developed by the Princeton Parallel Group. With it, one can easily create an SoC that has multiple CV6 cores and run a full-feature Operating System (OS) on a development board with an FPGA.

3.2.2 The Berkeley Out-of-Order RISC-V Processor

The Berkeley Out-of-Order RISC-V Processor (*BOOM* [33]) is a superscalar Out-of-Order (OoO) processor executing the RV64GC variant of the RISC-V ISA. BOOM was created at

the University of California, Berkeley in the Berkeley Architecture Research group. The CPU design is optimized to run on ASICs, although it can also be implemented on FPGAs. Its priority is to be a high-performance, synthesizable, and parametrizable core for architecture research. The current release, named *SonicBOOM*, has one of the best performances from the publicly available open-source RISC-V cores.

BOOM is a 10-stage CPU with the following stages: Fetch, Decode, Register Rename, Dispatch, Issue, Register Read, Execute, Memory, Writeback, and Commit. However, in most practical implementations, many of those stages are merged, generating seven stages altogether: Fetch, Decode/Rename, Rename/Dispatch, Issue/Register Read, Execute, Memory and Writeback. Since committing happens asynchronously, it is not counted as part of the “pipeline”. The load-store unit is optimized for the superscalar out-of-order architecture, and the data cache is organized into two dual-ported banks. At the front end, it is possible to customize the size of the L1 Instruction cache, the TLB, and the decode stage. Similarly to the CVA6, it is difficult or impossible to remove the cache from the core design and use the IOB-Cache instead.

This CPU design is written in Chisel [3] Hardware Description Language (HDL). The Constructing Hardware in a Scala Embedded Language (Chisel) allows for the production of synthesizable Verilog designs while using a high-level language to describe the hardware. Chisel is an adaptation of Scala [14] programming language, adding hardware construction primitives.

To build a System on a chip (SoC) with BOOM we would have to utilize the *Rocket Chip* [2] SoC generator from CHIPS Alliance. Since BOOM uses micro-architecture structures (TLBs, PTWs, etc) from that tool.

3.2.3 VexRiscv

The *VexRiscv* [16] CPU is a 32-bit Linux Capable RISC-V CPU written in the *SpinalHDL* [15]. The hardware description of this CPU is accomplished by utilizing a software-oriented approach. Similarly to Chisel, *SpinalHDL* is based on the Scala programming language.

VexRiscv is an in-order CPU with five “pipeline” stages. Many CPU plugins are optional, which add many functionalities to build a custom RISC-V CPU. The architecture design approach in this processor is unconventional, but it has its benefits: there are remarkably few fixed hardware components; Parts of the CPU can be swapped, turned on and turned off via the plugin system; without modifying any of the CPU sources, it is possible to add new functionalities/instructions easily; It permits the CPU arrangement to cover a significantly large spectrum of implementations, allowing the construction of an entirely parametrized CPU design. When the CPU is configured without plugins, it only includes the description of the five “pipeline” stages and their basic functionalities and nothing else. Everything else needs to be added to the CPU via plugins, including the program counter. VexRiscv can either be an application processor capable of

running a full-feature Operating System (OS) or a super simple microprocessor ideal for bare-bone applications depending on the way it is configured. Contrary to *BOOM*, *VexRiscv* does not need any external library. This makes it very easy to generate the synthesizable Verilog file from a *SpinalHDL* design.

There exists an open-source project that runs Linux with *VexRiscv*, *linux-on-litex-vexriscv* [11]. *LiteX* is used to create a System on a chip (SoC) around the *VexRiscv* core. *LiteX* SoC design and peripherals are written in *Migen* [6] another high level HDL. *Migen* unlike *SpinalHDL* and Chisel is based on Python 3.5. On account of the language describing its hardware and the way, the *linux-on-litex-vexriscv* project is structured it is very hard to understand how the system works, where the generated RTL is and how to add custom hardware. Furthermore, *linux-on-litex-vexriscv* uses FPGA specific hardware, making it impossible to port the system to ASIC.

Recently the developer behind *SpinalHDL* has also made public the *NaxRiscv* CPU. *NaxRiscv* is a CPU designed specifically to run a full-feature Operating System, like Linux. And just like *VexRiscv*, *NaxRiscv* uses *SpinalHDL* to describe its hardware. Although *NaxRiscv* seems like a very promising CPU it is still on its early stages. Consequently, it has a primitive interface which makes it complicated to implement on costume System on a chip (SoC).

3.3 Overall CPU comparison

In table 3.1 we can see a comparison of the CPUs that were presented in the previous sections, capable of running a full-feature Operating System (OS). All of the CPUs on the table are considered application processors. It can be observed that every CPU has a Memory Management Unit (MMU) and they all support U+S+M privilege mode. Furthermore, all of the CPUs hardware design have L1 Instruction Cache and L1 Data cache system integrated. This happens because to support atomic instructions it is easier to have direct access to the L1 Cache.

GNU/Linux is the combinations of *GNU* with the Linux kernel. The GNU Project developed a large part of the software that forms a complete Operating System (OS). Many “Linux” distributions make use of that software, a few examples would be *Debian*, *Ubuntu*, *openSUSE*, *Fedora*, and the list could go on. So a processor that supports the GNU/Linux feature is a CPU that is capable of running a distribution like *Ubuntu* or *Debian*. From the table, we can see that 32-bit RISC-V CPUs are the only ones not capable of running a *GNU/Linux* Operating System (OS).

	ARM	Andes Technology		SiFive		PULP platform	UC Berkeley	SpinalHDL	
	Cortex-A72	A25	AX25	U54	U74	CVA6	BOOM	VexRiscv	NaxRiscv
Architecture bit widths	64-bit	32-bit	64-bit	64-bit	64-bit	32/64-bit	64-bit	32-bit	64-bit
MMU	Y	Y	Y	Y	Y	Y	Y	Y	Y
FPU	Y	Y	Y	Y	Y	X	Y	X	Y
16-bit instructions	X	Y	Y	Y	Y	Y	Y	Y	X
Cache L1(I+D)	Y	Y	Y	Y	Y	Y	Y	Y	Y
Interrupt Controller	X	Y	Y	Y	Y	X	X	X	X
U+S+M Mode	N/A	Y	Y	Y	Y	Y	Y	Y	Y
GNU/Linux	Y	X	Y	Y	Y	Y	Y	X	Y
Open-Source	X	X	X	X	X	Y	Y	Y	Y

Table 3.1: CPU comparison table: Y means the CPU supports the feature; X means the CPU does not support the feature; N/A means the feature does not apply to the respective CPU.

4 | Hardware Developed

During the development of this thesis, there was both hardware and software developed. In this chapter, we are going to go through the hardware developed to build an appropriate System on a chip (SoC) capable of running a full-fledged Operating System (OS).

The *IOb-SoC* was used as a System on a chip (SoC) template. *IOb-SoC* has some features that make it ideal to develop this project SoC. Firstly, it is open-source hardware. This means there are no royalties and the source code is publicly available. Secondly, adding new peripherals is very easy and intuitive, as was previously seen in section 2.1. Thirdly, the *IOb-SoC* implements the interface with an internal (SRAM) and an external (DRAM) memory. When using an external memory the *IOb-SoC* instantiates an *iob-cache* system. Finally, the *IOb-SoC* implements a boot hardware unit that controls the first boot stage (also known as stage zero) that is executed after powering/resetting the system.

The hardware components that needed to be changed from *IOb-SoC* were the Central Processing Unit (CPU) and the Universal asynchronous receiver/transmitter (UART) peripheral. The CPU had to be changed because the previous CPU (*PicoRV32*) is not capable of running a full-feature Operating System. The UART had to be swapped since there were no compatible Linux drivers that worked with *iob-UART*. Besides swapping a few components from the chip new hardware had to be added. The additional hardware is the CLINT and the PLIC both compatible with RISC-V specifications. The CLINT was added to implement timer and software interrupts on the SoC. The PLIC was added to manage interrupts generated by other peripherals, for example from the UART. A sketch of the SoC developed can be seen in figure 4.1.

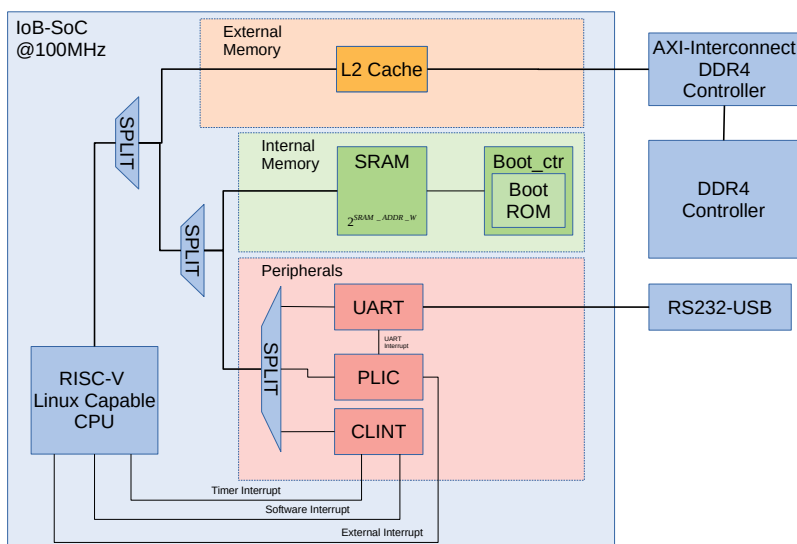


Figure 4.1: Developed SoC sketch.

Comparing figure 4.1 with the original design of *IOb-SoC* (figure 2.1) we can see that there were a few additional modifications. In the first place, it can be seen that the L1 Cache was removed. Since every application processor studied had an L1 cache built in, there was no need for the L1 *iob-cache*. Next, a *iob-split* was added to the *IOb-SoC*. Previously, there was a single *iob-split* for the data bus with three branches (the internal memory, the external memory and the peripheral bus). This meant that there were 2 selection bits, when '00' then the internal memory bus was active, when '01' it was the peripheral bus and when '10' it was the external memory. This caused a problem because when addressing the external memory if its size is bigger than 1GB the selection bits would be '11'. The demultiplexer (DEMUX) output selected by '11' is not connected anywhere, so this caused an internal hardware error. The solution was to include two *iob-split* modules each with two branches. The first would choose between the external memory and either the internal memory or peripheral bus. The second would choose between the internal memory and the peripheral bus. Another advantage of using this method is that now the selection bit's position does not vary depending on if we are using the DDR or not. This makes it easier to use external software that does not use the *iob-soc* Makefiles. Before the peripheral addressing on external software had to be changed every time the developer wanted to test with or without the external memory.

During this project, there was also an improvement on the *IOb-SoC* verification. This led to the creation of a top hardware module for the developed System on a chip.

4.1 Central Processing Unit

The CPU chosen to use in this project was *VexRiscv* [16]. The performance of the CPU is not a significant issue for this project. However, how the core was designed and developed highly influenced the CPU decision. The flexibility of the CPU design, meaning how easily the CPU can be adapted to take advantage of the other components in *IOb-SoC*, is an essential factor. Since the hardware and software developed in this project are open-source, the CPU implemented had to be open-source hardware. Moreover, knowing that the *IOb-SoC* signals are 32-bit wide, ideally, the selected CPU should support RV32IMAC to facilitate its integration with *IOb-SoC*. From the CPUs studied in chapter 3 *VexRiscv* looked like the more indicated.

Generating the RTL *verilog* file from the *SpinalHDL* hardware description is very simple. After cloning the *VexRiscv* GitHub repository the developer only has to run one command. As can be seen below in listing 4.1. On the *VexRiscv* repository there exist a couple of demo CPU configurations. The configurations can be directly used or configured to generate a custom *VexRiscv* core. There even already exists a demo configuration to generate a Linux-compatible core. Although in the developed hardware a custom core was implemented, the Linux demo configuration was used as a starting point. Unfortunately, the Linux Demo design is outdated and the instructions, commented on the hardware configuration file, to run a Linux simulation and test the core does not work.


```

1 git clone https://github.com/SpinalHDL/VexRiscv.git && \
2   cd VexRiscv && sbt "runMain vexriscv.demo.LinuxGen"

```

Listing 4.1: Generate *verilog* from *SpinalHDL*

The *VexRiscv* can be configured by adding and removing plugins. Plugins are hardware components described in *SpinalHDL* that can be reused in different designs by simply adding “*new Plugin_Name(...)*,” to the plugins list in the top CPU description file. The existing plugins are described in the *VexRiscv* repository on the “src/main/scala/vexriscv/plugin” directory. Looking at the available plugins it can be seen that there are two different plugins for the instruction bus and data bus. They are “*IBusSimplePlugin*”, “*IBusCachedPlugin*”, “*DBusSimplePlugin*” and “*DBusCachedPlugin*”. The difference is that the “cached” plugins have the L1 Cache integrated, while the simple plugins do not. An additional difference between the data cached and simple plugin is that, although the “*DBusCachedPlugin*” fully supports both, the “*DBusSimplePlugin*” supports only Load-Reserved (LR)/Store-Conditional (SC) but not Atomic Memory Operations (AMO) instructions. The “*DBusSimplePlugin*” could also be adapted to enable the full “A” extension. To do that it would be needed to have knowledge in *SpinalHDL*. Learning how to code in *SpinalHDL* would be very time-consuming and not in the scope of this project.

The first step on implementing the *VexRiscv* core on the *IOb-SoC* was making sure that it worked on “bare metal” applications. Meaning it had to be working with the application accessing the silicon chip directly without any intermediary like an Operating System (OS). This was done using the instruction and data “simple” plugins. The next step was to run the Linux kernel. To do so the instruction and data “simple” plugins had to be changed to the “cached” plugins. The missing support for Atomic Memory Operations (AMO) instructions was noticeable because the software would stop executing and enable an unknown instruction signal. It was possible to identify which instruction was causing the problem through the signal waves created during simulation.

The final *VexRiscv* core configuration file contained the needed plugins to run a minimal Operating System (OS) based on Linux. The plugins present were:

- The “*IBusCachedPlugin*” was added. With it, the address of the first instruction the CPU had to fetch was defined by setting the reset value of the Program counter (PC). Also, it was specified that the CPU had no branch predictor and that it supported compressed instructions. The decision to not use any branch predictor was because there seemed to be a compatibility problem between the most recent RISC-V toolchain and the branch predictor that are available in the *VexRiscv*. Since performance was not a concern in this project I choose to not use a branch predictor.
- The “*DBusCachedPlugin*” was added for the reason that it fully supported the atomic instructions.
- The “*DecoderSimplePlugin*” is used to decode the instructions.

- The “RegFilePlugin” implements the register file. These are the registers inside the CPU.
- The “IntAluPlugin” is used to calculate arithmetic and logic operations.
- The “SrcPlugin” is an auxiliary plugin for the plugins that contain arithmetic logic unit (ALU), Branch related hardware and Load/Store hardware logic.
- The “FullBarrelShifterPlugin” implements the shift instructions present in the RISC-V base Instruction set architecture (ISA).
- The “HazardSimplePlugin” is used by the core to determine where it needs to stall.
- The “MulPlugin” allows the core to execute multiplication instructions.
- The “MulDivIterativePlugin” could be used to add multiplication and division support to the core (RISC-V M ISA extension). In this case, it was used to add only division since the multiplication support was added by another plugin.
- The “CsrPlugin” is configured to fully support Linux. This plugin adds the needed Control and status register (CSR) to run a full feature OS.
- The “DebugPlugin” was deactivated in the used core. But it could be used to debug the CPU core if there existed a JTAG interface. Currently the *IOb-SoC* does not support it.
- The “BranchPlugin” allows the core to execute and make decisions on the jump instructions. This is part of the base Instruction set architecture (ISA)
- The “MmuPlugin” added support for the Memory Management Unit (MMU). Which is required to run a full feature OS.
- The “FpuPlugin” can add support for both the floats and doubles instruction extensions. In the core used this plugin was deactivated since to run a minimal OS there is little to no advantage of using this extension. Causing the FPU to only be adding unnecessary hardware logic.

After generating the Verilog file that describes a *VexRiscv* core I had to create a wrapper hardware module that adapted the *VexRiscv* core interface to the *IOb-SoC* internal bus.

4.1.1 VexRiscv Wrapper

The Verilog wrapper, which is called *iob_VexRiscv*, is instantiated by the *IOb-SoC* top SoC hardware module as the CPU component and instantiates the *VexRiscv* core Verilog module. The interface between the *IOb-SoC* hardware and the *VexRiscv* core is created by establishing a connection between the inputs and outputs from both sides.

The input signals of *iob_vexriscv* are the clock signal which is the system clock derivative from the development board where the SoC is implemented; the reset signal which is set to high

('1') when the system reboots and when the stage 0 bootloader finishes; the boot signal that has the value '1' while the stage 0 bootloader is executing, after it finishes the boot signal value drops to '0' at the same time the reset signal is set to high; the instruction bus response signal that is connected to "cpu_i_resp"; the data bus response signal that is connected to "cpu_d_resp"; the timer interrupt and software interrupt signals which are set to '1' or '0' by the CLINT unit; the external interrupt signal which is controlled by the PLIC unit. The output signals are the instruction bus request signal and the data bus request signal, which connect to the "cpu_i_req" and "cpu_d_req" respectively. The "cpu_i_resp", "cpu_d_resp", "cpu_i_req" and "cpu_d_req" signals were reviewed in section 2.1.

The input and output signals of the *VexRiscv* core can be seen in table 4.1. It can also be seen the signal's width and their equivalent signal in the *IOb-SoC* top hardware.

After understanding the inputs and outputs of each module it is easy to see which wires should be connected. But after connecting all the wires there were three problems. The first was the "strb" signal needed by *IOb-SoC* when writing data to memory which did not exist in the *VexRiscv* signals. The "strb" signal could be obtained in two different ways. One way would be through the "dbus_req_size" signal, the two less significant bits of the "dbus_req_address" and the "dbus_req_wr" signal. The other way was through the "dBus_cmd_payload_mask" signal and the "dBus_cmd_payload_wr" signal. The "DBusSimplePlugin", contrary to the "DBus-CachedPlugin", had no "dBus_cmd_payload_mask" signal that is why the first method was created. Accordingly for the first method, the "mask" signal had to be generated by the hardware logic expressed in equation 4.1.

$$\begin{cases} dbus_req_mask_aux = dbus_req_size[1]?4'hF : (dbus_req_size[0]?4'h3 : 4'h1) \\ dbus_req_mask = dbus_req_mask_aux << dbus_req_address[1 : 0] \end{cases} \quad (4.1)$$

Moreover, the "mask" signal indicated the active bytes when both read or write operations were occurring. On the other hand, the "strb" signal should only be active when a write operation is happening. Both methods logic expressions can be seen in equation 4.2. This implements a MUX where "dbus_req_wr" is the selection bit.

$$\begin{aligned} strb &= dbus_req_wr?dbus_req_mask : 4'h0 \\ &= dbus_req_wr?dBus_cmd_payload_mask : 4'h0 \end{aligned} \quad (4.2)$$

The second is that the *IOb-SoC* internal bus did not contain all the signals that were needed by the *VexRiscv* core. To successfully make the interface handshake with the *VexRiscv* core an instruction and data request "ready" signal had to be generated. The "ready" signal indicated that the SoC was ready to receive and accept a request from the CPU. To solve this problem a register that saved the value of the "cmd_valid", called "valid_reg", was created. This register

4 Hardware Developed

Port	Width	Direction	Description	IOb-SoC Port
dBus_cmd_valid	1	output	Indicates that the CPU is ready to make a data request.	cpu_d_req['valid(0)]
dBus_cmd_ready	1	input	Indicates that the SoC is ready to receive a data request.	N/A
dBus_cmd_payload_wr	1	output	Indicates that the CPU wants to write data to memory.	N/A
dBus_cmd_payload_uncached	1	output	Indicates if data is on L1 cache	Not used
dBus_cmd_payload_address	32	output	Used to address memory.	cpu_d_req['address(0,32)]
dBus_cmd_payload_data	32	output	Used to send data to memory.	cpu_d_req['wdata(0)]
dBus_cmd_payload_mask	4	output	Indicates which bytes in a word are accessed.	N/A
dBus_cmd_payload_size	2	output	\log_2 (number of bytes in the burst)	Not used
dBus_cmd_payload_last	1	output	Indicates when the last byte is transferred.	Not used
dBus_rsp_valid	1	input	Indicates that the SoC is ready to send a response.	cpu_d_resp['valid(0)]
dBus_rsp_payload_last	1	input	Indicates when the last byte is transferred.	Not used
dBus_rsp_payload_data	32	input	Receive data from memory.	cpu_d_resp['rdata(0)]
dBus_rsp_payload_error	1	input	Indicates existence of an error.	Not used
timerInterrupt	1	input	Indicate a Timer Interrupt.	timerInterrupt
externalInterrupt	1	input	Indicate an External Interrupt.	externalInterrupt
softwareInterrupt	1	input	Indicate a Software Interrupt.	softwareInterrupt
externalInterruptS	1	input	Indicate an External Interrupt at the Supervisor level.	Not used
iBus_cmd_valid	1	output	Indicates that the CPU is ready to make an instruction request.	cpu_i_req['valid(0)]
iBus_cmd_ready	1	input	Indicates that the SoC is ready to receive an instruction request.	N/A
iBus_cmd_payload_address	32	output	Used to address memory.	cpu_i_req['address(0,32)]
iBus_cmd_payload_size	2	output	\log_2 (number of bytes in the burst)	Not used
iBus_rsp_valid	1	input	Indicates that the SoC is ready to send a response.	cpu_i_resp['valid(0)]
iBus_rsp_payload_data	32	input	Receive an instruction from memory.	cpu_i_resp['rdata(0)]
iBus_rsp_payload_error	1	input	Indicates existence of an error.	Not used
clk	1	input	System clock signal.	clk
reset	1	input	CPU reset signal.	cpu_reset

Table 4.1: VexRiscv core inputs and outputs.

would be updated when either the “cmd_valid” or the “rsp_valid” signal were active. The “ready” signal should be high ('1') before accepting a request, and after it should be low ('0') while the response is not available. The initial approach to the values that the “cmd_ready” signal should assume can be seen in the truth table 4.2. This truth table was obtained by analyzing the simulation signals wave. The “N/A” values in the table mean that those situations never occurred.

The truth table can be transformed in a logic gates expression, which can be seen in equation 4.3.

valid_reg	cmd_valid	rsp_valid	cmd_ready
0	0	0	0
0	0	1	N/A
0	1	0	1
0	1	1	N/A
1	0	0	0
1	0	1	1
1	1	0	0
1	1	1	1

Table 4.2: First try at identifying the rules the cmd_ready should follow.

$$(valid_reg \cdot rsp_valid) + (valid \cdot \overline{valid_reg} \cdot \overline{rsp_valid}) \quad (4.3)$$

But this approach had an issue. The logic expression depended on the “cmd_valid” signal which was generated inside the *VexRiscv* core. This could generate a bigger complication since the combinatorial circuit that generates the “cmd_valid” is unknown and might generate an infinite hardware loop. From better analyzing the signal behavior it was noticed that when “valid_reg”, “cmd_valid” and “rsp_valid” are low ('0') the value of “cmd_ready” is irrelevant. It could be concluded since when the “valid_reg”, “cmd_valid” and “rsp_valid” are low happens the “cmd_ready” signal was not being used by the *IOb-SoC*. The truth table can then be simplified to table 4.3.

valid_reg	rsp_valid	req_ready
0	0	1
0	1	N/A
1	0	0
1	1	1

Table 4.3: Simplified truth table.

Which can be seen as a simple XOR logic gate. The equation 4.4 show the hardware logic expression implement.

$$(valid_reg \cdot rsp_valid) + (\overline{valid_reg} \cdot \overline{rsp_valid}) = valid_reg \odot rsp_valid \quad (4.4)$$

The last problem was that after accepting an instruction or data request the values of the “address”, “data” and “mask” signals could change inside the *VexRiscv* core. This changes would pass through *ioB_VexRiscv* and reflect in the rest of *IOb-SoC* hardware. Which caused the *ioB-cache* and peripherals to not function currently. This problem was solved by creating registers that saved the value of the “address”, “data” and “strb” signals when the request

was accepted. The register values would then only change when the response was already received.

To finalize, the CPU should be able to run firmware from both the internal and external memory. When the stage 0 bootloader is running the Most Significant bit (msb) of the instruction fetched address had to be forced to '0'. This would force the CPU to fetch instructions from the boot hardware unit. When defined that the firmware had to run from the external memory (RUN_EXTMEM=1) the first instruction fetched should be at address `0x80000000`. To achieve this requirement the Most Significant bit (msb), when RUN_EXTMEM=1 was defined as the negated value of the boot signal. When RUN_EXTMEM=0 the msb was forced to always be '0' since there is no need to access the external memory. On the data request bus, it should also be taken into account that the msb had to be '0' when the CPU wanted to access the peripherals.

4.2 UART 16550

The approach taken in this project was to adapt an existing open-source Universal asynchronous receiver/transmitter (UART) core that is supported by the Linux kernel. The other option was to create a Linux driver compatible with *ioB-UART* and compile the kernel with it. The chosen approach seemed more adequate and a simpler solution.

Since the developed chip is supposed to be open-source the UART core should also be open-source hardware. The core used was a *UART16550* [9] that has been made available by *freecores* on *github*. This UART was written in Verilog, although it was an older version of Verilog it is still synthesizable by modern tools and easy to understand. The *UART16550* core used implements a wishbone interface to interact with the System on a chip (SoC). Similarly to what was done with the CPU, I had to create a wrapper to adapt the core to the *IOB-SoC*.

4.2.1 UART 16550 Wrapper

The wishbone interface is established in the top hardware module from the used *UART16550* core. In table 4.4, which can be obtained from the open-source *UART16550* core documentation, it can be seen the wishbone interface signals. The wishbone specification determines that there needs to be a master and a slave. In this case when the CPU sends a request signal for the *UART16550* peripheral the master is the CPU and the slave is the *UART16550*.

The interface between the *IOB-SoC* and the *UART16550* top hardware is established by the verilog module that acts as a *UART16550* wrapper. The wrapper module is called "ioB_uart16550". The "ioB_uart16550" hardware component has to generate the missing signal that are needed

Port	Width	Direction	Description
CLK	1	Input	Block's clock input
WB_RST_I	1	Input	Asynchronous Reset
WB_ADDR_I	5 or 3	Input	Used for register selection
WB_SEL_I	4	Input	Select signal
WB_DAT_I	32 or 8	Input	Data input
WB_DAT_O	32 or 8	Output	Data output
WB_WE_I	1	Input	Write or read cycle selection
WB_STB_I	1	Input	Specifies transfer cycle
WB_CYC_I	1	Input	A bus cycle is in progress
WB_ACK_O	1	Output	Acknowledge of a transfer

Table 4.4: WISHBONE interface signals.

by the wishbone interface but do not exist in *IOb-SoC*. In table 4.5 the connection between the *IOb-SoC* and the *UART16550* wishbone interface can be analyzed. The missing signal are the “WB_SEL_I”, the “WB_WE_I” and the “WB_STB_I”. The select signal is similar to the “strb” signal but it should exist during write and read operations. This signal can be obtained through the two Less Significant bit (lsb) of the address signal. And since registers are addressed byte by byte only one bit at a time will be set to high on the select signal. The write signal is set to high whenever the CPU wants to write to the *UART16550* registers. This can be perceived through the *IOb-SoC* “strb” signal. If any of the “strb” bits are enabled the “WB_WE_I” signal should be high ('1'). The transfer cycle should happen (i.e. “WB_STB_I” signal should be set to '1') when the *UART16550* has accepted a request but still has not issued the response. Furthermore, the *UART16550* has an interrupt output pin that is connected to the SoC “uartInterrupt”. The SoC “uartInterrupt” is then passed to the PLIC unit.

UART16550 Wishbone	IOb-UART16550	IOb-SoC
CLK	clk	clk
WB_RST_I	rst	reset
WB_ADDR_I	address['UART_ADDR_WIDTH-1:0]	slaves_req['address('UART16550,32)]
WB_SEL_I	1<<address[1:0]	N/A
WB_DAT_I	wdata	slaves_req['wdata('UART16550)]
WB_DAT_O	rdata	slaves_resp['rdata('UART16550)]
WB_WE_I	wstrb	N/A
WB_STB_I	valid&(~ready)	N/A
WB_CYC_I	valid	slaves_req['valid('UART16550)]
WB_ACK_O	ready	slaves_resp['ready('UART16550)]

Table 4.5: *UART16550* interface with *IOb-SoC*.

Finally, the “iob_uart16550” also has to pass the interface between the *UART16550* core and the RS232 connector. That is why it implements the: txd output to transmit data through serial;

the rxd input to receive data through serial; the cts input which indicates that the destination is ready to receive a transmission sent by the UART; the rts output which indicates that the UART is ready to receive a transmission from the sender. Those pins are connected to the development board RS232 connector.

4.3 CLINT Unit

The CLINT was the only hardware component that was developed from scratch. Even though, there already exist open-source Core-local Interrupt Controller (CLINT) hardware modules. For example, the CLINT used with the *CVA6* core and developed by the *PULP platform*. The problem with this CLINT module is that it is written in system-Verilog and uses packages and definitions from the *CVA6* core. The CLINT core is a simple hardware component. Considering that it only needs a few registers and signals to work, as it was studied in section 2.2. The interface with *IOb-SoC* would have to be developed independently of the CLINT core used. The best solution was to fully create the CLINT hardware unit.

The inputs and outputs of the CLINT unit can be seen in table 4.6. “N_CORES” is the number of CPU core that are used in the SoC. In the System on a chip (SoC) developed there is only one core. But the CLINT is built with a multi-core system in mind. Each core has its timer and software interrupt. As such, the “mtip” and “msip” registers width is the number of core, “N_CORES”.

Port	Width	Direction	Description	IOb-SoC Port
clk	1	input	System clock	clk
rst	1	input	System reset	reset
rt_clk	1	input	Real-time clock	rtc
valid	1	input	Indicates that the CPU is ready to make a data request.	slaves_req[‘valid(‘CLINT)]
address	32	input	Register address.	slaves_req[‘address(‘CLINT,16)]
wdata	32	input	Data to write to register.	slaves_req[‘wdata(‘CLINT)]
wstrb	4	input	Used to generate a "write" signal.	slaves_req[‘wstrb(‘CLINT)]
rdata	32	output	Data read from register.	slaves_resp[‘rdata(‘CLINT)]
ready	1	output	Indicates that the CLINT is ready to send a response.	slaves_resp[‘ready(‘CLINT)]
mtip	N_CORES	output	Raise a timer interrupt in a core.	timerInterrupt
msip	N_CORES	output	Raise a software interrupt in a core.	softwareInterrupt

Table 4.6: CLINT interface with *IOb-SoC*.

The “rt_clk” signal, although it is connected to “rtc” wire, is not available in *IOb-SoC* since the development board does not have any arithmetic logic unit (rtc) crystal connected. The rtc frequency is commonly 32.768 kHz, because it is a power of 2 (2^{15}) value. To get a precise 1-second period (1 Hz frequency) would only be needed a 15 stage binary counter. With a rtc the CLINT unit had to detect the rising edge of the “rt_clk” signal. The rising edge was

detected by taking samples of the “rt_clk” signal at the system clock frequency. This is possible because the rtc is slower than the system clock. If in the future the CLINT unit developed is implemented in a system with a rtc the hardware design can be easily adapted to it. Since the logic was already developed.

In the development boards used there is no rtc so an alternative had to be found. The method implemented was simpler than the logic if there existed a rtc. The system clock operates at 100 MHz so a counter was added to the design and each time that counter reached 999 the timer register would increment. Like this, a rtc working at 100 kHz was simulated in the CLINT hardware.

Concluding, the needed registers to develop a CLINT unit in accordance with the RISC-V specifications were implemented. And the hardware logic needed to trigger the timer and software interrupts were described.

4.4 PLIC Unit

The PLIC is not essential to run a full feature Operating System (OS) on a System on a chip (SoC). Since the PLIC is used to drive interrupts generated by other peripherals to the core, in this project the only peripheral connected to the PLIC is the *UART16550*. For the growth of the SoC, the PLIC unit is handy and a requirement to have. Some peripherals that might be added in the future also use the PLIC hardware, for example, the *ethernet* controller can be used to wake a core from low power mode.

Since the PLIC hardware unit is more complex than the CLINT, it was decided that an open-source PLIC core would be adapted to the *IOb-SoC*. There were three available cores. The PLIC developed by *lowRISC* [13] is written in their own variation of *System Verilog*, consequently it is difficult to adapt and test in simulation. The PLIC used with the *CVA6* core, developed by the *PULP platform* [17], started as a fork of an older version of the *lowRISC*. This PLIC unit is written in traditional *System Verilog*. I tried to adapt the PLIC unit from *PULP platform* but there were many incompatibilities with the *IOb-SoC*. Finally, the PLIC developed by *RoaLogic* [19] was the hardware unit used as a starting point. The *RoaLogic* PLIC is also written in *System Verilog* and implements an *apb4* or *ahb3lite* interface with the SoC. The biggest advantage of the *RoaLogic* hardware is that the PLIC relevant components for the *IOb-SoC*, the PLIC registers and core, are well separated from the modules that create the interface with the SoC and instantiate them. To integrate the a PLIC unit on the *IOb-SoC* I had to develop a PLIC wrapper.

4.4.1 PLIC wrapper

The PLIC wrapper had to create the interface with the *IOb-SoC* internal buses. Furthermore, it needed to instantiate the PLIC registers and core hardware modules.

The PLIC hardware module was named “*job_plic*”. When instantiating the “*job_plic*” there are various parameters the developer can customize that influence the hardware. The number of interrupt sources defines how many peripherals interrupt signals are connected to the PLIC. In this project, the number of interrupt sources is one (the *UART16550*). The number of interrupt targets is normally equal to the number of CPU cores. The PLIC will send an external interrupt signal to each target. The number of priority levels in the developed SoC is 8 and each interrupt source has a register that holds the value for its priority. The maximum number of ‘pending’ events is the maximum number of peripheral interrupts the PLIC will register while waiting for the CPU to solve the current interrupt. There is also a parameter that defines if ‘threshold’ is implemented. The last configurable parameter determines whether the ‘configuration’ register is implemented. The changes in the parameters influence the number of registers and wires the most. The hardware logic does not change.

The “*job_plic*” interface with the *IOb-SoC* can be seen in table 4.7. The value of “SOURCES” is 8. It was decided to be this value because although right now the SoC has only one source in the future it might have up to 8 sources. The value of “TARGETS” is 1 since there is only one core to inform of an existing interrupt.

Port	Width	Direction	Description	IOb-SoC Port
clk	1	input	System clock	clk
rst	1	input	System reset	reset
valid	1	input	CPU is sending a request.	slaves_req[‘valid(‘PLIC)]
address	32	input	Register address.	slaves_req[‘address(‘PLIC, 16)]
wdata	32	input	Data to write to register.	slaves_req[‘wdata(‘PLIC)]
wstrb	4	input	Used to generate a write and a read enable signal.	slaves_req[‘wstrb(‘PLIC)]
rdata	32	output	Data read from register.	slaves_resp[‘rdata(‘PLIC)]
ready	1	output	PLIC is ready to send a response.	slaves_resp[‘ready(‘PLIC)]
src	SOURCES	input	Peripherals Interrupts	{uartInterrupt, {7{1'b0}}}
irq	TARGETS	output	Inform targets of existing interrupt.	{externalInterrupt}

Table 4.7: PLIC interface with *IOb-SoC*.

Furthermore, the PLIC response signal is always ready one clock cycle after a request is received. Additionally, the PLIC has write enable and read enable signals. The write enable signal can be generated from equation 4.5. The read enable is obtained similarly but when all bits of the “strb” signal are ‘0’, equation 4.6.

$$valid \cdot (wstrb[3] + wstrb[2] + wstrb[1] + wstrb[0]) \quad (4.5)$$

$$valid \cdot \overline{(wstrb[3] + wstrb[2] + wstrb[1] + wstrb[0])} \quad (4.6)$$

5 | Software Developed

Simultaneously with the development of the hardware logic components there was software developed. During this project first it was developed a software for communication with the *IOb-SoC* through serial. Second it was developed a new hardware simulation software. Thirdly, firmware that could test and run interrupt routines was developed. Fourth, the needed software/firmware to execute an Operating System was adapted and compiled for the SoC developed. Finally, there were multiple Makefiles written that facilitate the user interaction and further development.

Some of the software developed was not mandatory to get a full-fledged Operating System to work with the System on a chip designed. The complementary software was matured because it facilitates the project development. Furthermore, the additional software allows the *IOb-SoC* platform to support more features.

In this chapter the software developed will be analyzed. The new *Python Console* and the new simulation system were already implemented on the upstream *IOb-SoC* template. This software is already being used by other developers. Moreover, it has been improved by the *IObundle* developers using it.

5.1 *Python Console*

The *Console* is a program that runs on the user computer and communicates with the board where the *IOb-SoC* is implemented using an *RS-232* connection. Initially the *IOb-SoC* had a *Console* written in *C* programming language. One of the first tasks developed was the translation of the *Console* program to *Python*.

The *C Console* makes use of a set of functions on a independent file that were written to read/write to the serial port. The *Python* program uses the *PySerial* library, which provides ready-made communication functions like those in the original *C* code. Using *PySerial* is better because the community regularly maintains and updates *PySerial*. *PySerial* provides additional features, is less prone to have bugs, and the communication is more trustworthy comparing with the *C* functions.

One of the reasons to translate the *Console* program was to integrate an existing Ethernet controller already written in *Python*. *Python* can easily exploit feature like files, sockets and other Operating System (OS) functionalities.

The *Python Console* program can be used in two different modes: locally working with simulators, or communicating with a board running *IOb-SoC*. The program mode can be choose when calling the *Console* through adding “-L” or “-local” to the invoking arguments. This is an alteration to the original *Console* program. The *C Console* could only works with the FPGA

board. When the *Console* is run in board mode a physical implementation of *IOb-SoC* runs on the board and communicates with *Console* through a *RS-232* serial connection. If the *Console* is called with the “-L” or “-local” augment it will communicate with the simulator. The communication with the hardware simulation is identical to the one with the board. They exchange the same messages. When communicating with the simulator the *Console* uses files to send and receive data from the *IOb-SoC* hardware simulation. The *Console* program when starting creates two empty files in the simulation directory. The “cns12soc” is used to send messages from the *Console* to the SoC. The “soc2cns1” is used by the *Console* to receive messages from the SoC. Both files only contain one byte at a time. The fact if the files are empty or not is used to synchronize the simulation with the *Console*. After reading from one of the files the simulation or the *Console* program has to empty the respective file.

The way the code is structured is very similar to how it was on the *C Console* program. It starts by defining the parameters that influence messages identifiers and serial communication (for example, the number of bits per byte, the parity and the number of stop bits). When the program enters its primary function, it starts a loop where it waits for an available byte to read, either from the serial port or the file, depending on the mode the *Console* was called. After receiving the byte from the SoC, it computes what type of message it is. The program exits successfully if the byte received is an End of Transmission (EOT = 04, ASCII value in hexadecimal). If the byte received is an Enquiry (ENQ = 05, ASCII value in hexadecimal), the program checks if it was the first time it received an enquiry. If it was, it could have one of either behavior: if the program was called with an argument equivalent to “-f”, meaning that there is a firmware that should be uploaded to the *IOb-SoC*, the *Console* sends a *Send a file request* message (FRX = 08, value in hexadecimal) to the *IOb-SoC*; if there is not a firmware file to send then the *Console* responds to the *IOb-SoC* with an Acknowledgement (ACK = 06, ASCII value in hexadecimal). If the *Console* receives a Receive a file request message (FTX = 07, value in hexadecimal), it will run a function that will receive any file sent from the *IOb-SoC* to the computer and save it under the directory where it is running. If the *Console* receives a *Send a file request* message, it will run a function that will send any file requested from the *IOb-SoC* to it. Any other byte received will be printed onto the stdout.

The FRX and the FTX bytes are specific to the *IOb-SoC* platform software. This could cause a problem when using external software that does not attribute the same meaning to their respective values. To solve this problem a meaning for the Device Control 1 (DC1 = 11, ASCII value in hexadecimal) byte was created. When receiving a DC1 byte the *Console* deactivates all platform specific meanings for the respective bytes. This means that after receiving a DC1 byte stops associating the value 0x07, 0x08, 0x11 to FTX, FRX and DC1 respectively.

An example of how to call the *Console* to communicate with the simulation and send the firmware to the *IOb-SoC* when it starts would be 5.1. The “&” at the end means that the *Console* program is going to run in the background. Allowing to run other programs while the *Console* executes.

```
CONSOLE_CMD=$(CONSOLE_DIRECTORY)/console -L -f &
```

Listing 5.1: Call *Console* program

5.2 IOb-SoC Simulation

In order to support the new *Console* simulation mode a new *IOb-SoC* verification mechanism had to be developed. Verification is one an important topic when developing hardware. A correct and precise verification saves time since the hardware does not have to be synthesized and flashed to an FPGA every time an SoC designer wants to test a new feature.

The original simulation **How it worked...** The responses given to the *IOb-SoC* were performed with a emulation of the *Console* program written in Verilog. Every time the *Console* was updated, the simulation *Console* also had to be updated. Hence, the idea was to create a test-bench that allowed the simulator to interact with the *Console* program. The new simulation now has the advantage of mostly using the same *Console* program as when the *IOb-SoC* is implemented in a FPGA.

5.2.1 UUT Top Hardware module

This top module creates a *verilog* wrapper of the Unit Under Test (UUT) that allows it to interact with the different hardware logic simulators. Even tho this is written hardware logic it is never implemented as real hardware. This module is only used in simulation as software.

The top module file is an adaptation of the previous *verilog* file used on the Icarus simulation. Previously, the verilog top file would interact directly with the *Console*. Similarly to the new hardware top module an UART module is instantiated and a serial connection with the *IOb-SoC* is simulated. Although in this project the *IOb-UART* in the System on a chip (SoC) was swapped for the *UART16550* in the

5.2.2 Verilator Testbench

Verilator transforms the Verilog HDL designs into a C++ program that can be executed after being compiled. Using C++ to create a testbench permits to execute the converted hardware program simulating the hardware initially described in Verilog. While also allowing to easily make use of system calls. The testbench needed to run with Verilator is similar to the testbench in Verilog used with Icarus.

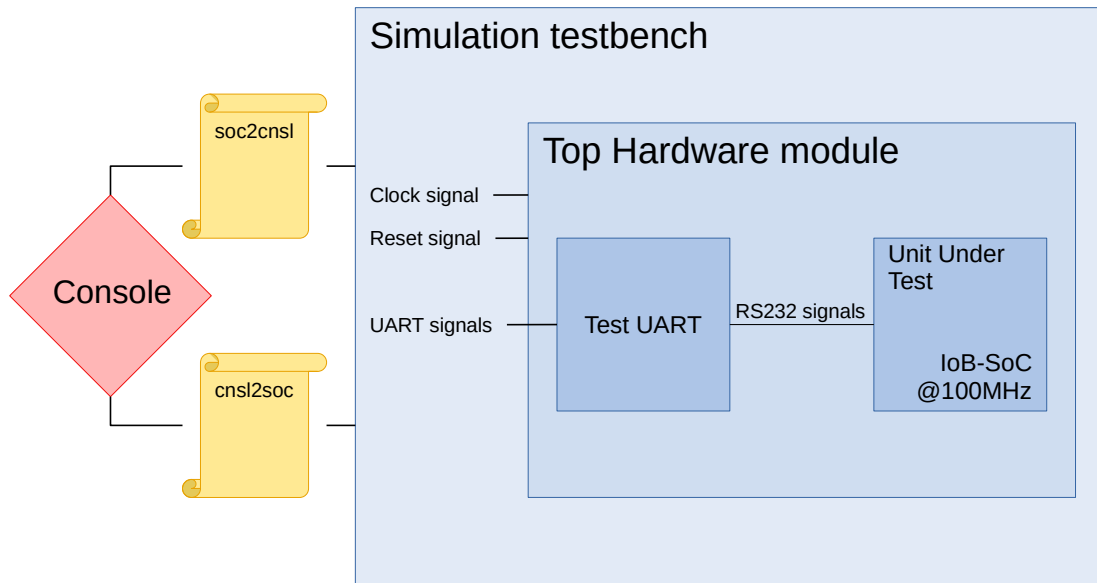


Figure 5.1: Simulated hardware interfaces.

5.3 Interrupt Routine

During the development of the CLINT hardware unit, no firmware used the CLINT features. The CLINT enables the support for time or software-related interrupts. Therefore, to test the CLINT hardware it was needed to create a simulation testbench. Moreover, to understand how interrupts are used in code and handled by the CPU a bare-metal firmware that uses interrupts was created.

To generate a time-related interrupt the software has to write to the “MTIMECMP” register. The “MTIMECMP” register address is the “MTIMECMP_BASE” address, which is 0x4000, plus 0x08 times the core id. The core id is the “TARGET”, CPU core, which is supposed to receive the interrupt notification. In this project since there is only one core, the core id is 0. Furthermore, when writing firmware to run on the SoC it is needed to consider the CLINT peripheral base address and add it to the “MTIMECMP” register address. To make a timer interrupt trigger after 10 seconds the firmware has to do more than just write to the “MTIMECMP”. First, the firmware has to read the current time. The current time can be obtained from the “MTIME” register. Although not directly since the “MTIME” register increments with the CLINT designed frequency, in this case 100MHz. To convert the value in “MTIME” register to seconds we know that $seconds = \frac{* (MTIME)}{frequency}$. The “MTIME” register address is obtained similarly to the “MTIMECMP” register but instead of the “MTIMECMP_BASE”, the “MTIME_BASE”, which is 0xbff8, is used. After reading the current time, the value that needs to be stored in “MTIMECMP” register is calculated. The value is calculated by adding the time to wait before the interrupt is triggered to the current time. The value calculated can then be stored in the “MTIMECMP” register. When the “MTIME” register value is equal or greater than the “MTIMECMP” register

value the timer interrupt is enabled. The pseudo-code to set up the timer interrupt can be seen in the code snippet 5.2.

```

1  #define MTIMECMP_BASE 0x4000
2  #define MTIME_BASE 0xbff8
3  #define FREQ 100000000
4  void set_up_mtip(time_sec){
5      long long aux_value = 0; // 64-bit integer
6      int core_id = 0;
7      aux_value = *(MTIME_BASE+8*core_id)
8      aux_value = aux_value + time_sec*FREQ;
9      *(MTIMECMP_BASE+8*core_id) = aux_value;
10 }

```

Listing 5.2: Set Up Timer Interrupt.

The core id could have been read from the CSR which saves its value. The RISC-V instruction that does so is “csrr %0, mhartid”. An example of a C code integration would be code snippet 5.3.

```

1  static inline uint_32_t csr_read_mhartid(void) {
2      uint_32_t value;
3      __asm__ volatile ("csrr    %0, mhartid"
4                          : "=r" (value) /* output : register */
5                          : /* input : none */
6                          : /* clobbers: none */);
7      return value;
8  }

```

Listing 5.3: Read core id from CSR.

One of the software interrupt usages is to synchronize various cores in a system. When dividing the workload between cores there might be a time when core 1 has to synchronize with core 0. To do this core 1 would wait until core 0 generates a software interrupt targeting core 1. In this project, there is only one core so this situation does not occur. But applications can run concurrently with multi-threading. One application could wait until a software interrupt is triggered. This software interrupt could be triggered by another application running in core 0 and targeting core 0. To generate a software-related interrupt the software has to write to the “MSIP” register. The “MSIP” register address is the “MSIP_BASE” address, which is 0x00, plus 0x04 times the core id. When running the firmware adding the CLINT peripheral base address to the “MSIP” register can not be overlooked. The “MSIP” register can only be written through external sources. The CLINT hardware cannot change it internally. The pseudo-code to set up the software interrupt can be seen in the code snippet 5.4.

```

1  #define MSIP_BASE 0x00
2  void set_up_msip(){
3      int core_id = csr_read_mhartid();
4      *(MSIP_BASE+4*core_id) = 1;
5  }

```

Listing 5.4: Set Up Software Interrupt.

After enabling an interrupt the CLINT sends a hardware notification. The interrupt notification has to be handled by the rest of the hardware. The CLINT testbench and the bare-metal firmware handle the interrupt notification differently.

5.3.1 CLINT simulation

To test the CLINT hardware unit a simulation testbench was built in Verilog and another in C++ programming language. This allows the developers to test the correctness of the hardware component without connecting it to the rest of the SoC.

Both the testbench in Verilog and the testbench in C++ had similar behaviours. First, they would set up a timer interrupt. The timer interrupt was set to trigger in $0.2 * 10^{-6}$ seconds. Taking into account that the simulation is slow this was a reasonable time. After the timer interrupt is triggered the simulation receives its notification and proceeds to handle the interrupt. When receiving a timer interrupt the simulation sets up the software interrupt. In the next clock cycle, the testbench is notified of an existing software interrupt. It then proceeds to disable the timer and the software interrupt. To disable the timer interrupt the “MTIMECMP” register has to be set to its maximum value, which is 0xFFFFFFFFFFFFFFFF (i.e. all 64 bits are '1'). To disable the software interrupt the “MSIP” register had to be set to 0.

If the interrupts work correctly there will be a message in the terminal indicating their correctness. After testing that both interrupts work as expected the testbench can finish successfully. The simulation will always end $1 * 10^{-6}$ seconds after starting.

5.3.2 Bare-metal firmware

Once the CLINT unit was developed it had to be tested while integrated into the SoC. Even though it was known that the CLINT generated the interrupts through the simulation testbench. To test the CLINT in the SoC firmware that took advantage of the timer and/or software interrupts had to be developed.

Since the CLINT hardware developed is compatible with RISC-V, any firmware compatible with RISC-V that took advantage of interrupts should work. With this in mind, the open-source bare-metal firmware made available by *Five EmbedDev* [8], an embedded RISC-V blog, would be taken advantage of. The firmware could not all be used directly in *IOb-SoC*. Some functions that interact with the timer in the CLINT were adapted to the developed hardware and used. Furthermore, a file that implemented functions that read/write to the Control and status register is used as a library.

The developed firmware.c, similarly to the original *IOb-SoC* firmware, starts by initializing the UART and the CLINT hardware. Then it will disable all global interrupts. To disable the global interrupts the “mstatus” CSR, “mie” CSR and the “mcause” CSR are cleared. After the timer

interrupt is set similarly to the code in 5.2. When an interrupt occurs the program counter has to jump to a specific function. The memory address of that function is saved in the “mtvec” CSR. Succeeding the timer interrupt set up, the respective interrupt bit can be set to ‘1’ (i.e. enable the timer interrupt) in the “mie” CSR. Following this the global interrupts can be enabled again by setting the needed bits in “mstatus” CSR to ‘1’. the bits that need to be set to ‘1’ are the ones corresponding to the machine interrupts (MSTATUS_MIE_BIT_MASK). Finally, the program can wait for an interrupt to happen with the “wfi” instruction.

When an interrupt occurs the interrupt handler function is called. The interrupt handler will read what was the reason the interrupt was generated from the “mcause” CSR. After knowing the cause the interrupt handler will act accordingly to how it was programmed. The developed firmware informs the user that a timer interrupt occurred and it sets the “MTIMECMP” register to a higher number. The developed firmware finishes after receiving the first interrupt.

It is important to note that besides the firmware.c some alterations had to be made in the IOb-SoC firmware.S file. In the firmware.S the global pointer had to be set. The global pointer is similar to the stack pointer. The difference is while the stack pointer point to the memory location where functions variables will be stored, the global pointer points to the memory location where global variables are stored. When setting the global pointer it is critical to write the ‘norelax’ option in the Assembly code. Without “.option norelax” the global pointer will be loaded relative to the global pointer.

5.4 IOb-SoC Linux OS integration

Talk about the noncanonical.py

5.4.1 Bootloaders

IOb-SoC Linux Stage 0 Bootloader

OpenSBI

5.4.2 Device Tree

5.4.3 Linux kernel

5.4.4 Root File system

5.5 Makefiles

6 | Project Results

6.1 System Running "Hello World!"

6.2 Interrupt Routines

6.3 Run/Boot Linux Performance

6.4 FPGA Resources Consumption

7 | Conclusions

7.1 Achievements

7.2 Contributed Repositories

- **iob-soc**
- **iob-soc-vexriscv**
- **iob-lib**
- **iob-vexriscv**
- **iob-uart16550**
- **iob-clint**
- **iob-plic**

7.3 Future Work

- ethernet
- development board with IOb-SoC

A | Annex

A.1 Annex 1

Annex 1

A.2 Annex 2

Annex 2

Bibliography

- [1] Arm Holdings (arm). *The Cortex-A72 processor specifications*. <<https://developer.arm.com/Processors/Cortex-A72>>.
- [2] Krste Asanovic et al. "The rocket chip generator". In: *EECS Department, University of California, Berkeley, Tech. Rep. UCB/EECS-2016-174* (2016).
- [3] Jonathan Bachrach et al. "Chisel: constructing hardware in a scala embedded language". In: *DAC Design automation conference 2012*. IEEE. 2012, pp. 1212–1221.
- [4] Jonathan Balkind et al. "OpenPiton: An Open Source Manycore Research Framework". In: *Proceedings of the Twenty-First International Conference on Architectural Support for Programming Languages and Operating Systems*. ASPLOS '16. Atlanta, Georgia, USA: ACM, 2016, pp. 217–232. ISBN: 978-1-4503-4091-5. DOI: 10.1145/2872362.2872414. <<http://doi.acm.org/10.1145/2872362.2872414>>.
- [5] Seeed Studio BeagleBoard.org and StarFive. *BeagleV, The First Affordable RISC-V Computer Designed to Run Linux*. <<https://beagleboard.org/static/beagleV/beagleV.html>>.
- [6] Sebastien Bourdeauducq. "Migen Manual". In: *Release 0. X 32* (2012).
- [7] darklife. *DarkRISCV*. <<https://github.com/darklife/darkriscv>>.
- [8] Five EmbedDev. *riscv-scratchpad - baremetal-startup-c*. <<https://github.com/five-embeddev/riscv-scratchpad/tree/master/baremetal-startup-c/src>>.
- [9] Jacob Gorban. "Uart ip core specification". In: *Architecture 15* (2002), p. 1.
- [10] IObundle Lda. *IOb-SoC*. <<https://github.com/IObundle/iob-soc>>.
- [11] LiteX. *linux-on-litex-vexriscv*. <<https://github.com/litex-hub/linux-on-litex-vexriscv>>.
- [12] lowRISC. *Ibex RISC-V Core*. <<https://github.com/lowRISC/ibex>>.
- [13] lowRISC. *OpenTitan - PLIC*. <https://github.com/lowRISC/opentitan/tree/master/hw/ip_templates/rv_plic>.
- [14] Martin Odersky et al. *The Scala language specification*. 2004.
- [15] C Papon. "SpinalHDL: An alternative hardware description language". In: *FOSDEM*. 2017.
- [16] C Papon. *VexRiscv*. <<https://github.com/SpinalHDL/VexRiscv>>.
- [17] pulp-platform. *RISC-V Platform-Level Interrupt Controller*. <https://github.com/pulp-platform/rv_plic/tree/master>.
- [18] riscv. *RISC-V Platform Specification*. <<https://github.com/riscv/riscv-platform-specs/blob/main/riscv-platform-spec.pdf>>.

- [19] RoaLogic. *AHB-Lite Platform-Level Interrupt Controller (PLIC)*. <<https://github.com/RoaLogic/plic>>.
- [20] SiFive. *HiFive Unleashed*. <<https://www.sifive.com/boards/hifive-unleashed>>.
- [21] SiFive. *HiFive Unmatched*. <<https://www.sifive.com/boards/hifive-unmatched>>.
- [22] SiFive. *U54*. <<https://www.sifive.com/cores/u54>>.
- [23] SiFive. *U54-MC*. <<https://www.sifive.com/cores/u54-mc>>.
- [24] SiFive. *U74*. <<https://www.sifive.com/cores/u74>>.
- [25] SiFive. *U74-MC*. <<https://www.sifive.com/cores/u74-mc>>.
- [26] stnolting. *The NEORV32 RISC-V Processor*. <<https://github.com/stnolting/neorv32>>.
- [27] Andes Technology. *ADP-XC7K160/410, FPGA Based Development Platform*. <<http://www.andestech.com/en/products-solutions/andeshape-platforms/adp-xc7k160-410/>>.
- [28] Andes Technology. *AE350 Platform*. <<http://www.andestech.com/en/products-solutions/andeshape-platforms/ae350-axi-based-platform-pre-integrated-with-n25f-nx25f-a25-ax25/>>.
- [29] Andes Technology. *AndesCore™ A25, Compact High-Speed 32-bit CPU for Real-time and Linux Applications*. <<http://www.andestech.com/tw/%E7%94%A2%E5%93%81%E8%88%87%E8%A7%A3%E6%B1%BA%E6%96%B9%E6%A1%88/andescore-processors/riscv-a25/>>.
- [30] Andes Technology. *AndesCore™ AX25, Compact High-Speed 64-bit CPU for Real-time and Linux Applications*. <<http://www.andestech.com/tw/%E7%94%A2%E5%93%81%E8%88%87%E8%A7%A3%E6%B1%BA%E6%96%B9%E6%A1%88/andescore-processors/riscv-ax25/>>.
- [31] YosysHQ. *PicoRV32 - A Size-Optimized RISC-V CPU*. 2015. <<https://github.com/YosysHQ/picorv32>>.
- [32] F. Zaruba and L. Benini. "The Cost of Application-Class Processing: Energy and Performance Analysis of a Linux-Ready 1.7-GHz 64-Bit RISC-V Core in 22-nm FDSOI Technology". In: *IEEE Transactions on Very Large Scale Integration (VLSI) Systems* 27.11 (Oct. 2019), pp. 2629–2640. ISSN: 1557-9999. DOI: 10.1109/TVLSI.2019.2926114.
- [33] Jerry Zhao et al. "SonicBOOM: The 3rd Generation Berkeley Out-of-Order Machine". In: (May 2020).